

# Exercícios Cap 05

Rodrigo Giannotti

## Capítulo 05

### Inicialização

```
library(tidyverse)
library(magrittr) # mais pipes, como %<>%
library(lubridate) # melhor manejo de datas
```

Para o capítulo 5 também utilizaremos a biblioteca de voos de NYC

```
library(nycflights13)
# ?flights
# View(flights)
head(flights)
```

```
## # A tibble: 6 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>    <int>         <int>
## 1  2013     1     1     517             515         2      830             819
## 2  2013     1     1     533             529         4      850             830
## 3  2013     1     1     542             540         2      923             850
## 4  2013     1     1     544             545        -1     1004            1022
## 5  2013     1     1     554             600        -6      812             837
## 6  2013     1     1     554             558        -4      740             728
## # ... with 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>
```

## Exercícios

### 5.2 filter()

#### 5.2.1

Find all flights that:

**a** Had an arrival delay of two or more hours

```
flights %>% filter(  
  arr_delay >= 120  
)
```

```
## # A tibble: 10,200 x 19  
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time  
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>  
## 1  2013     1     1     811             630        101    1047             830  
## 2  2013     1     1     848             1835       853    1001             1950  
## 3  2013     1     1     957             733        144    1056             853  
## 4  2013     1     1    1114             900        134    1447             1222  
## 5  2013     1     1    1505             1310       115    1638             1431  
## 6  2013     1     1    1525             1340       105    1831             1626  
## 7  2013     1     1    1549             1445         64    1912             1656  
## 8  2013     1     1    1558             1359       119    1718             1515  
## 9  2013     1     1    1732             1630         62    2028             1825  
## 10 2013     1     1    1803             1620       103    2008             1750  
## # ... with 10,190 more rows, and 11 more variables: arr_delay <dbl>,  
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

**b** Flew to Houston (IAH or HOU)

```
flights %>% filter(  
  dest %in% c("IAH", "HOU")  
)
```

```
## # A tibble: 9,313 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time  
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>  
## 1  2013     1     1     517           515         2      830           819  
## 2  2013     1     1     533           529         4      850           830  
## 3  2013     1     1     623           627        -4      933           932  
## 4  2013     1     1     728           732        -4     1041          1038  
## 5  2013     1     1     739           739         0     1104          1038  
## 6  2013     1     1     908           908         0     1228          1219  
## 7  2013     1     1    1028          1026         2     1350          1339  
## 8  2013     1     1    1044          1045        -1     1352          1351  
## 9  2013     1     1    1114           900        134     1447          1222  
## 10 2013     1     1    1205          1200         5     1503          1505
```

```
## # ... with 9,303 more rows, and 11 more variables: arr_delay <dbl>,
```

```
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
```

```
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

c Were operated by United, American or Delta

```
flights %>% filter(  
  carrier %in% c("UA", "AA", "DL")  
)
```

```
## # A tibble: 139,504 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time  
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>  
## 1  2013     1     1     517           515         2      830           819  
## 2  2013     1     1     533           529         4      850           830  
## 3  2013     1     1     542           540         2      923           850  
## 4  2013     1     1     554           600        -6      812           837  
## 5  2013     1     1     554           558        -4      740           728  
## 6  2013     1     1     558           600        -2      753           745  
## 7  2013     1     1     558           600        -2      924           917  
## 8  2013     1     1     558           600        -2      923           937  
## 9  2013     1     1     559           600        -1      941           910  
## 10 2013     1     1     559           600        -1      854           902  
## # ... with 139,494 more rows, and 11 more variables: arr_delay <dbl>,  
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

d Departed in the summer (July, August and September)

```
summer <- c(7:9)
flights %>% filter(
  month %in% summer
)
```

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     7     1       1           2029          212     236           2359
## 2  2013     7     1       2           2359           3     344           344
## 3  2013     7     1      29           2245          104     151             1
## 4  2013     7     1      43           2130          193     322            14
## 5  2013     7     1      44           2150          174     300            100
## 6  2013     7     1      46           2051          235     304           2358
## 7  2013     7     1      48           2001          287     308           2305
## 8  2013     7     1      58           2155          183     335             43
## 9  2013     7     1     100           2146          194     327             30
## 10 2013     7     1     100           2245          135     337            135
## # ... with 86,316 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

e Arrived more than two hours late, but didn't leave late

```
flights %>% filter(
  arr_delay >= 120 & dep_delay <= 0
)
```

```
## # A tibble: 29 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>
## 1  2013     1    27    1419           1420        -1    1754           1550
## 2  2013    10     7    1350           1350         0    1736           1526
## 3  2013    10     7    1357           1359        -2    1858           1654
## 4  2013    10    16     657             700        -3    1258           1056
## 5  2013    11     1     658             700        -2    1329           1015
## 6  2013     3    18    1844           1847        -3     39           2219
## 7  2013     4    17    1635           1640        -5    2049           1845
## 8  2013     4    18     558             600        -2    1149            850
## 9  2013     4    18     655             700        -5    1213            950
## 10 2013     5    22    1827           1830        -3    2217           2010
## # ... with 19 more rows, and 11 more variables: arr_delay <dbl>, carrier <chr>,
## #   flight <int>, tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>,
## #   distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

f Were delayed by at least an hour, but made up over 30 minutes in flight

```
flights %>% filter(  
  dep_delay >= 60 & (dep_delay - arr_delay) >= 30  
)
```

```
## # A tibble: 2,074 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time  
##   <int> <int> <int>   <int>         <int>      <dbl>   <int>         <int>  
## 1  2013     1     1    1716           1545        91    2140           2039  
## 2  2013     1     1    2205           1720       285     46           2040  
## 3  2013     1     1    2326           2130       116    131            18  
## 4  2013     1     3    1503           1221       162   1803          1555  
## 5  2013     1     3    1821           1530       171   2131          1910  
## 6  2013     1     3    1839           1700        99   2056          1950  
## 7  2013     1     3    1850           1745        65   2148          2120  
## 8  2013     1     3    1923           1815        68   2036          1958  
## 9  2013     1     3    1941           1759       102   2246          2139  
## 10 2013     1     3    1950           1845        65   2228          2227
```

```
## # ... with 2,064 more rows, and 11 more variables: arr_delay <dbl>,
```

```
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
```

```
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

g Departed between 00:00 and 6:00 (inclusive)

```
flights %>% filter(  
  dep_time <= 600 | dep_time == 2400  
)
```

## # A tibble: 9,373 x 19

##	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
##	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
## 1	2013	1	1	517	515	2	830	819
## 2	2013	1	1	533	529	4	850	830
## 3	2013	1	1	542	540	2	923	850
## 4	2013	1	1	544	545	-1	1004	1022
## 5	2013	1	1	554	600	-6	812	837
## 6	2013	1	1	554	558	-4	740	728
## 7	2013	1	1	555	600	-5	913	854
## 8	2013	1	1	557	600	-3	709	723
## 9	2013	1	1	557	600	-3	838	846
## 10	2013	1	1	558	600	-2	753	745

## # ... with 9,363 more rows, and 11 more variables: arr\_delay <dbl>,

## # carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,

## # air\_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time\_hour <dtm>



### 5.2.2

Another useful dplyr filtering helper is `between()`. What does it do? Can you use it to simplify the code needed to answer the previous challenges?

```
# ?between
```

Como dito na ajuda, “This is a shortcut for `x >= left & x <= right`” ou seja, é uma maneira de testar se valores dentro de um vetor estão dentro de dois limites.

Isso só seria útil para simplificar a questão dos meses do verão

```
flights %>% filter(
  between(month, 7, 9)
)
```

```
## # A tibble: 86,326 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     7     1       1           2029           212     236           2359
## 2  2013     7     1       2           2359            3     344           344
## 3  2013     7     1      29           2245          104     151            1
## 4  2013     7     1      43           2130          193     322            14
## 5  2013     7     1      44           2150          174     300           100
## 6  2013     7     1      46           2051          235     304           2358
## 7  2013     7     1      48           2001          287     308           2305
## 8  2013     7     1      58           2155          183     335            43
## 9  2013     7     1     100           2146          194     327            30
## 10 2013     7     1     100           2245          135     337           135
## # ... with 86,316 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

### 5.2.3

How many flights have a missing `dep_time`? What other variables are missing? What might these rows represent?

```
flights %>% filter(
  is.na(dep_time)
)
```

```
## # A tibble: 8,255 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>       <dbl>   <int>         <int>
## 1  2013     1     1     NA             1630         NA       NA             1815
## 2  2013     1     1     NA             1935         NA       NA             2240
## 3  2013     1     1     NA             1500         NA       NA             1825
## 4  2013     1     1     NA              600         NA       NA              901
## 5  2013     1     2     NA             1540         NA       NA             1747
## 6  2013     1     2     NA             1620         NA       NA             1746
## 7  2013     1     2     NA             1355         NA       NA             1459
## 8  2013     1     2     NA             1420         NA       NA             1644
## 9  2013     1     2     NA             1321         NA       NA             1536
##10  2013     1     2     NA             1545         NA       NA             1910
## # ... with 8,245 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Podemos ver que os voos com `dep_time` vazio apresentam outras colunas vazias, como `dep_delay`; `arr_time`; `arr_delay`; e `air_time`

Isso muito provavelmente indica voos que foram cancelados

#### 5.2.4

Why is  $NA \wedge 0$  not missing? Why is  $NA \mid TRUE$  not missing? Why is  $FALSE \& NA$  not missing? Can you figure out the general rule? ( $NA * 0$  is a tricky counterexample!)

```
# help(`^`)  
# help(`/`)
```

Como podemos ver no texto de ajuda “ $1 \wedge y$  and  $y \wedge 0$  are 1, always.”, dessa forma o operador nem passa pela etapa de avaliar o NA, simplesmente retornando o resultado.

Similarmente, “NA is a valid logical object. Where a component of x or y is NA, the result will be NA if the outcome is ambiguous. In other words  $NA \& TRUE$  evaluates to NA, but  $NA \& FALSE$  evaluates to FALSE. See the examples below.”

Logo como sempre  $(x \mid TRUE)$  retornaria TRUE e  $(x \& FALSE)$  retornaria FALSE independentemente dos valores de x, logo retornam-se os valores lógicos.

Isso só ocorre quando o computador está explicitamente tomando a decisão de não avaliar a expressão como um todo, devido à um de seus lados. não existe tão decisão para  $NA * 0$ , por exemplo, logo o resultado esperado é NA.

```
NA * 0
```

```
## [1] NA
```

## 5.3 arrange()

### 5.3.1

How could you use `arrange()` to sort all missing values to the start? (Hint: use `is.na()`.)

```
flights %>% arrange(  
  desc(  
    is.na(dep_time)  
  ))
```

```
## # A tibble: 336,776 x 19
```

```
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time  
##   <int> <int> <int>   <int>         <int>         <dbl>    <int>         <int>  
## 1  2013     1     1     NA             1630           NA        NA             1815  
## 2  2013     1     1     NA             1935           NA        NA             2240  
## 3  2013     1     1     NA             1500           NA        NA             1825  
## 4  2013     1     1     NA              600           NA        NA              901  
## 5  2013     1     2     NA             1540           NA        NA             1747  
## 6  2013     1     2     NA             1620           NA        NA             1746  
## 7  2013     1     2     NA             1355           NA        NA             1459  
## 8  2013     1     2     NA             1420           NA        NA             1644  
## 9  2013     1     2     NA             1321           NA        NA             1536  
## 10 2013     1     2     NA             1545           NA        NA             1910  
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,  
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

### 5.3.2

Sort flights to find the most delayed flights. Find the flights that left earliest.

```
flights %>% arrange(
  desc(
    (dep_delay + arr_delay)
  )) # maior atraso somado entre saída e chegada
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     9     641             900         1301    1242         1530
## 2  2013     6    15    1432            1935         1137    1607         2120
## 3  2013     1    10    1121            1635         1126    1239         1810
## 4  2013     9    20    1139            1845         1014    1457         2210
## 5  2013     7    22     845            1600         1005    1044         1815
## 6  2013     4    10    1100            1900          960    1342         2211
## 7  2013     3    17    2321             810          911     135         1020
## 8  2013     7    22    2257             759          898     121         1026
## 9  2013    12     5     756            1700          896    1058         2020
## 10 2013     5     3    1133            2055          878    1250         2215
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights %>% arrange(
  dep_delay
) # menor atraso de saída
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013    12     7    2040            2123          -43     40         2352
## 2  2013     2     3    2022            2055          -33    2240         2338
## 3  2013    11    10    1408            1440          -32    1549         1559
## 4  2013     1    11    1900            1930          -30    2233         2243
## 5  2013     1    29    1703            1730          -27    1947         1957
## 6  2013     8     9     729             755          -26    1002          955
## 7  2013    10    23    1907            1932          -25    2143         2143
## 8  2013     3    30    2030            2055          -25    2213         2250
## 9  2013     3     2    1431            1455          -24    1601         1631
## 10 2013     5     5     934             958          -24    1225         1309
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

### 5.3.3

Sort flights to find the fastest (highest speed) flights.

```
flights %>% arrange(
  desc(
    distance / air_time
  )
)
```

## # A tibble: 336,776 x 19

##	year	month	day	dep_time	sched_dep_time	dep_delay	arr_time	sched_arr_time
##	<int>	<int>	<int>	<int>	<int>	<dbl>	<int>	<int>
## 1	2013	5	25	1709	1700	9	1923	1937
## 2	2013	7	2	1558	1513	45	1745	1719
## 3	2013	5	13	2040	2025	15	2225	2226
## 4	2013	3	23	1914	1910	4	2045	2043
## 5	2013	1	12	1559	1600	-1	1849	1917
## 6	2013	11	17	650	655	-5	1059	1150
## 7	2013	2	21	2355	2358	-3	412	438
## 8	2013	11	17	759	800	-1	1212	1255
## 9	2013	11	16	2003	1925	38	17	36
## 10	2013	11	16	2349	2359	-10	402	440

## # ... with 336,766 more rows, and 11 more variables: arr\_delay <dbl>,  
## # carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
## # air\_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time\_hour <dtm>

### 5.3.4

Which flights traveled the longest? Which traveled the shortest?

```
flights %>% arrange(
  desc(
    distance
  )) # voos mais longos
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1     857             900          -3    1516         1530
## 2  2013     1     2     909             900           9    1525         1530
## 3  2013     1     3     914             900          14    1504         1530
## 4  2013     1     4     900             900           0    1516         1530
## 5  2013     1     5     858             900          -2    1519         1530
## 6  2013     1     6    1019             900          79    1558         1530
## 7  2013     1     7    1042             900         102    1620         1530
## 8  2013     1     8     901             900           1    1504         1530
## 9  2013     1     9     641             900        1301    1242         1530
##10  2013     1    10     859             900          -1    1449         1530
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

```
flights %>% arrange(
  distance
) # voos mais curtos
```

```
## # A tibble: 336,776 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     7    27      NA             106          NA      NA           245
## 2  2013     1     3    2127             2129          -2    2222         2224
## 3  2013     1     4    1240             1200          40    1333         1306
## 4  2013     1     4    1829             1615         134    1937         1721
## 5  2013     1     4    2128             2129          -1    2218         2224
## 6  2013     1     5    1155             1200          -5    1241         1306
## 7  2013     1     6    2125             2129          -4    2224         2224
## 8  2013     1     7    2124             2129          -5    2212         2224
## 9  2013     1     8    2127             2130          -3    2304         2225
##10  2013     1     9    2126             2129          -3    2217         2224
## # ... with 336,766 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

## 5.4 select()

### 5.4.1

Brainstorm as many ways as possible to select `dep_time`, `dep_delay`, `arr_time`, and `arr_delay` from `flights`.

Por mais que uma partida de regex golf sempre tenha seu valor nostalgico a função `starts_with()` resolve o problema com mais simplicidade

```
flights %>% select(  
  starts_with("dep_")  
  |  
  starts_with("arr_")  
)
```

```
## # A tibble: 336,776 x 4  
##   dep_time dep_delay arr_time arr_delay  
##   <int>     <dbl>   <int>     <dbl>  
## 1      517         2     830         11  
## 2      533         4     850         20  
## 3      542         2     923         33  
## 4      544        -1    1004        -18  
## 5      554        -6     812        -25  
## 6      554        -4     740         12  
## 7      555        -5     913         19  
## 8      557        -3     709        -14  
## 9      557        -3     838         -8  
## 10     558        -2     753          8  
## # ... with 336,766 more rows
```



### 5.4.2

What happens if you include the name of a variable multiple times in a `select()` call?

```
flights %>% select(  
  dep_time, dep_time, dep_time, arr_time, dep_time  
)
```

```
## # A tibble: 336,776 x 2  
##   dep_time arr_time  
##   <int>    <int>  
## 1      517      830  
## 2      533      850  
## 3      542      923  
## 4      544     1004  
## 5      554      812  
## 6      554      740  
## 7      555      913  
## 8      557      709  
## 9      557      838  
## 10     558      753  
## # ... with 336,766 more rows
```

Somente uma copia dessa coluna chega ao resultado final

### 5.4.3

What does the `one_of()` function do? Why might it be helpful in conjunction with this vector?

```
vars <- c(
  "year", "month", "day", "dep_delay", "arr_delay"
)
```

```
# ?tidyselect::one_of
```

Como podemos ver o próprio tidyverse sugere o uso das mais precisas `all_of()` ou `any_off()`, que servem para - em conjunto com o comando `select()` - selecionar variáveis com nomes dentro de listas. o comando `all_of` retorna erro se algum dos nomes da lista não for encontrado como nome de coluna enquanto o `any_off` ignora as colunas que não forem encontradas. No caso ambos devem retornar o mesmo dataframe, visto que todas as colunas da lista existem.

```
flights %>% select(
  any_of(
    vars
  )
)
```

```
## # A tibble: 336,776 x 5
##   year month   day dep_delay arr_delay
##   <int> <int> <int>     <dbl>     <dbl>
## 1  2013     1     1         2         11
## 2  2013     1     1         4         20
## 3  2013     1     1         2         33
## 4  2013     1     1        -1        -18
## 5  2013     1     1        -6        -25
## 6  2013     1     1        -4         12
## 7  2013     1     1        -5         19
## 8  2013     1     1        -3        -14
## 9  2013     1     1        -3         -8
## 10 2013     1     1        -2          8
## # ... with 336,766 more rows
```

#### 5.4.4

Does the result of running the following code surprise you? How do the select helpers deal with case by default? How can you change that default?

```
select(flights, contains("TIME"))

## # A tibble: 336,776 x 6
##   dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
##   <int>      <int>      <int>      <int>      <dbl> <dtm>
## 1      517          515      830          819      227 2013-01-01 05:00:00
## 2      533          529      850          830      227 2013-01-01 05:00:00
## 3      542          540      923          850      160 2013-01-01 05:00:00
## 4      544          545     1004         1022      183 2013-01-01 05:00:00
## 5      554          600      812          837      116 2013-01-01 06:00:00
## 6      554          558      740          728      150 2013-01-01 05:00:00
## 7      555          600      913          854      158 2013-01-01 06:00:00
## 8      557          600      709          723       53 2013-01-01 06:00:00
## 9      557          600      838          846      140 2013-01-01 06:00:00
## 10     558          600      753          745      138 2013-01-01 06:00:00
## # ... with 336,766 more rows

# ?select
```

Isso não surpreende, mas poderia ser interessante esperar que o helper “contains()” considerasse caixa baixa ou alta. Para mudar seu comportamento para que passe a considerar isso basta usar o argumento `ignore.case = FALSE`

```
flights %>% select(
  contains(
    "TIME", ignore.case = F
  )
)
```

```
## # A tibble: 336,776 x 0
```

## 5.5 mutate()

### 5.5.1

Currently `dep_time` and `sched_dep_time` are convenient to look at, but hard to compute with because they're not really continuous numbers. Convert them to a more convenient representation of number of minutes since midnight.

```
(flights2 <- flights %>% mutate(  
  dep_time_mins = (dep_time %/% 100) * 60 + dep_time %% 100,  
  sched_dep_time_mins = (sched_dep_time %/% 100) * 60 + sched_dep_time %% 100  
)  
  
## # A tibble: 336,776 x 21  
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time  
##   <int> <int> <int>   <int>         <int>      <dbl>    <int>         <int>  
## 1  2013     1     1     517           515         2      830           819  
## 2  2013     1     1     533           529         4      850           830  
## 3  2013     1     1     542           540         2      923           850  
## 4  2013     1     1     544           545        -1     1004          1022  
## 5  2013     1     1     554           600        -6      812           837  
## 6  2013     1     1     554           558        -4      740           728  
## 7  2013     1     1     555           600        -5      913           854  
## 8  2013     1     1     557           600        -3      709           723  
## 9  2013     1     1     557           600        -3      838           846  
## 10 2013     1     1     558           600        -2      753           745  
## # ... with 336,766 more rows, and 13 more variables: arr_delay <dbl>,  
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,  
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>,  
## #   dep_time_mins <dbl>, sched_dep_time_mins <dbl>
```

### 5.5.2

Compare `air_time` with `arr_time - dep_time`. What do you expect to see? What do you see? What do you need to do to fix it?

Primeiro vamos dar o mesmo tratamento que demos para os horários de saída para os horários de chegada


```
flights2 %<>% mutate(  
  sched_arr_time_mins = (sched_arr_time %/% 100) * 60 + sched_arr_time %% 100,  
  arr_time_mins = (arr_time %/% 100) * 60 + arr_time %% 100  
)
```

Com isso vamos criar uma nova coluna, que compara o horário efetivo de saída com o de chegada, tomando alguns cuidados, especialmente com voos que saem em uma data e chegam em outra. Supoe-se que nenhum voo voa por mais de 24h nesse caso (e uma rápida olhada para a coluna `air_time` confirma isso)

```
flights2 %<>% mutate(  
  arr_dep_time_diff = ifelse(  
    arr_time_mins >= dep_time_mins, # se o voo chegou no dia seguinte a conta arr_time - dep_time daria  
    arr_time_mins - dep_time_mins, # caso não hajam problemas  
    arr_time_mins - dep_time_mins + 24*60) # adicionando o numero de minutos em um dia caso o voo tenha  
)
```

Agora vamos comparar a diferença entre horário de saída com a coluna `air_time`

```
flights2 %>%  
  na.omit() %>%  
  mutate(  
    comparacao =  
      air_time - arr_dep_time_diff  
) %>%  
  select(comparacao) %>%  
  boxplot()
```



```
flights2 %>%
  na.omit() %>%
  mutate(
    comparacao =
      air_time - arr_dep_time_diff
  ) %>%
  select(comparacao) %>%
  summary()
```

```
##   comparacao
##   Min.    :-204.00
##   1st Qu.: -30.00
##   Median : -19.00
##   Mean   :  13.38
##   3rd Qu.:  38.00
##   Max.    : 345.00
```

Esperaria-se que essa comparação fosse sempre igual a 0, visto que se espera que o tempo de voo se iguale à diferença de hora de saída e hora de chegada, o problema é que neste caso estamos ignorando o tempo de taxi e que existam imperfeições de registro. (poderíamos ter questões de fuso horário, mas ?flights nos informa que todos os horários estão na timezone de NYC)

### 5.5.3

Compare `dep_time`, `sched_dep_time`, and `dep_delay`. How would you expect those three numbers to be related?

```
flights2 %>% mutate(  
  comparison = ifelse(  
    dep_time_mins >= (sched_dep_time_mins + dep_delay),  
    dep_time_mins - (sched_dep_time_mins + dep_delay),  
    dep_time_mins - (sched_dep_time_mins + dep_delay) + 24*60  
  )) %>%  
  select(comparison) %>%  
  summary()
```

```
##      comparison  
##  Min.      :0  
## 1st Qu.:0  
##  Median :0  
##   Mean  :0  
## 3rd Qu.:0  
##   Max.  :0  
##  NA's   :8255
```

Esperaria-se que essa comparação retornasse valor nulo, o que confirmamos com o comando `summary()`

#### 5.5.4

Find the 10 most delayed flights using a ranking function. How do you want to handle ties? Carefully read the documentation for `min_rank()`.

```
flights %>%
  mutate(
    rank_delay = min_rank( desc(
      arr_delay + dep_delay
    )) %>%
    arrange(rank_delay) %>%
    filter(rank_delay <= 10)

## # A tibble: 10 x 20
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     9     641             900         1301    1242         1530
## 2  2013     6    15    1432            1935         1137    1607         2120
## 3  2013     1    10    1121            1635         1126    1239         1810
## 4  2013     9    20    1139            1845         1014    1457         2210
## 5  2013     7    22     845            1600         1005    1044         1815
## 6  2013     4    10    1100            1900          960    1342         2211
## 7  2013     3    17    2321             810          911     135         1020
## 8  2013     7    22    2257             759          898     121         1026
## 9  2013    12     5     756            1700          896    1058         2020
## 10 2013     5     3    1133            2055          878    1250         2215
## # ... with 12 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
## #   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
## #   hour <dbl>, minute <dbl>, time_hour <dtm>, rank_delay <int>

# ?rank
```

Por sorte não temos empates usando essa formula, mas se tivessemos por usar a forma `min_rank` teriamos que empates teriam o mesmo rank, este sendo o minimo entre eles.



### 5.5.5

What does `1:3 + 1:10` return? Why?

```
1:3 + 1:10
```

```
## Warning in 1:3 + 1:10: longer object length is not a multiple of shorter object
```

```
## length
```

```
## [1]  2  4  6  5  7  9  8 10 12 11
```

a lógica é `1+1 2+2 3+3` e então, como a lista menor acaba temos `1+4 2+5 3+6 1+7 2+8 3+9 1+10` E um aviso informando que a lista maior não tem uma quantidade de itens múltipla da quantidade de itens na lista menor.

### 5.5.6

What trigonometric functions does R provide?

```
# ?sin
```

O texto de ajuda das funções trigonométricas lista as seguintes: "These functions give the obvious trigonometric functions. They respectively compute the cosine, sine, tangent, arc-cosine, arc-sine, arc-tangent, and the two-argument arc-tangent."

## 5.6 summarise()

Para essa seção em alguns exercícios utilizaremos uma tabela com os voos que foram cancelados.

```
not_cancelled <- flights %>%  
  filter(!is.na(dep_delay), !is.na(arr_delay))
```

### 5.6.1

Brainstorm at least 5 different ways to assess the typical delay characteristics of a group of flights. Consider the following scenarios:

- a A flight is 15 minutes early 50% of the time, and 15 minutes late 50% of the time.
- b A flight is always 10 minutes late.
- c A flight is 30 minutes early 50% of the time, and 30 minutes late 50% of the time.
- d 99% of the time a flight is on time. 1% of the time it's 2 hours late.
- e Which is more important: arrival delay or departure delay?

### 5.6.2

Come up with another approach that will give you the same output as `not_cancelled %>% count(dest)` and `not_cancelled %>% count(tailnum, wt = distance)` (without using `count()`).

```
identical(  
  not_cancelled %>% count(dest),  
  not_cancelled %>%  
    group_by(dest) %>%  
    summarise(n = n(), .groups = 'drop')  
)
```

```
## [1] TRUE
```

```
identical(  
  not_cancelled %>% count(tailnum, wt = distance),  
  not_cancelled %>%  
    group_by(tailnum) %>%  
    summarise(n = sum(distance), .groups = 'drop')  
)
```

```
## [1] TRUE
```

### 5.6.3

Our definition of cancelled flights (`is.na(dep_delay) | is.na(arr_delay)`) is slightly suboptimal. Why? Which is the most important column?

Primeiro, vamos olhar para a coluna com a maior quantidade de NAs

```
flights %>% filter(is.na(arr_delay))
```

```
## # A tibble: 9,430 x 19
##   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
##   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>
## 1  2013     1     1    1525           1530         -5    1934           1805
## 2  2013     1     1    1528           1459         29    2002           1647
## 3  2013     1     1    1740           1745         -5    2158           2020
## 4  2013     1     1    1807           1738         29    2251           2103
## 5  2013     1     1    1939           1840         59      29           2151
## 6  2013     1     1    1952           1930         22    2358           2207
## 7  2013     1     1    2016           1930         46      NA           2220
## 8  2013     1     1      NA           1630         NA      NA           1815
## 9  2013     1     1      NA           1935         NA      NA           2240
## 10 2013     1     1      NA           1500         NA      NA           1825
## # ... with 9,420 more rows, and 11 more variables: arr_delay <dbl>,
## #   carrier <chr>, flight <int>, tailnum <chr>, origin <chr>, dest <chr>,
## #   air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
```

Encontramos alguns voos com horário de saída, horário de chegada, mas nem tempo de voo nem atraso de chegada.


É difícil saber como considerar se estes voos foram ou não cancelados, mas me parece que o mais correto seria considerar um voo cancelado como aquele que não decolou, e portanto, tem `dep_time == NA`

#### 5.6.4

Look at the number of cancelled flights per day. Is there a pattern? Is the proportion of cancelled flights related to the average delay?


É de se imaginar que se existe algum padrão de cancelamentos ou atrasos estes se dariam em datas especiais ou em dias da semana mais comuns. Para isso usando a coluna `time_hour` e o pacote `lubridate` podemos criar colunas para ter mais insights sobre esses padrões.

```
flights3 <- flights %>%  
  mutate(dia_semana = lubridate::wday(time_hour))  
  
flights3 %>%  
  group_by(dia_semana) %>%  
  summarise(soma_cancelados = sum(is.na(dep_time)), .groups = "drop") %>%  
  ggplot(mapping = aes(x = dia_semana, y = soma_cancelados)) +  
  geom_col(fill = "lightblue") # voos cancelados por dia da semana
```



Primeiro podemos ver que muito menos voos foram cancelados durante o fim de semana, mas será que existe uma correlação com o número total de voos nesses dias?

```
flights3 %>%  
  group_by(dia_semana) %>%  
  summarise(n_voos = n(), .groups = "drop") %>%  
  ggplot(mapping = aes(x = dia_semana, y = n_voos)) +  
  geom_col(fill = "lightblue") # voos totais por dia da semana
```



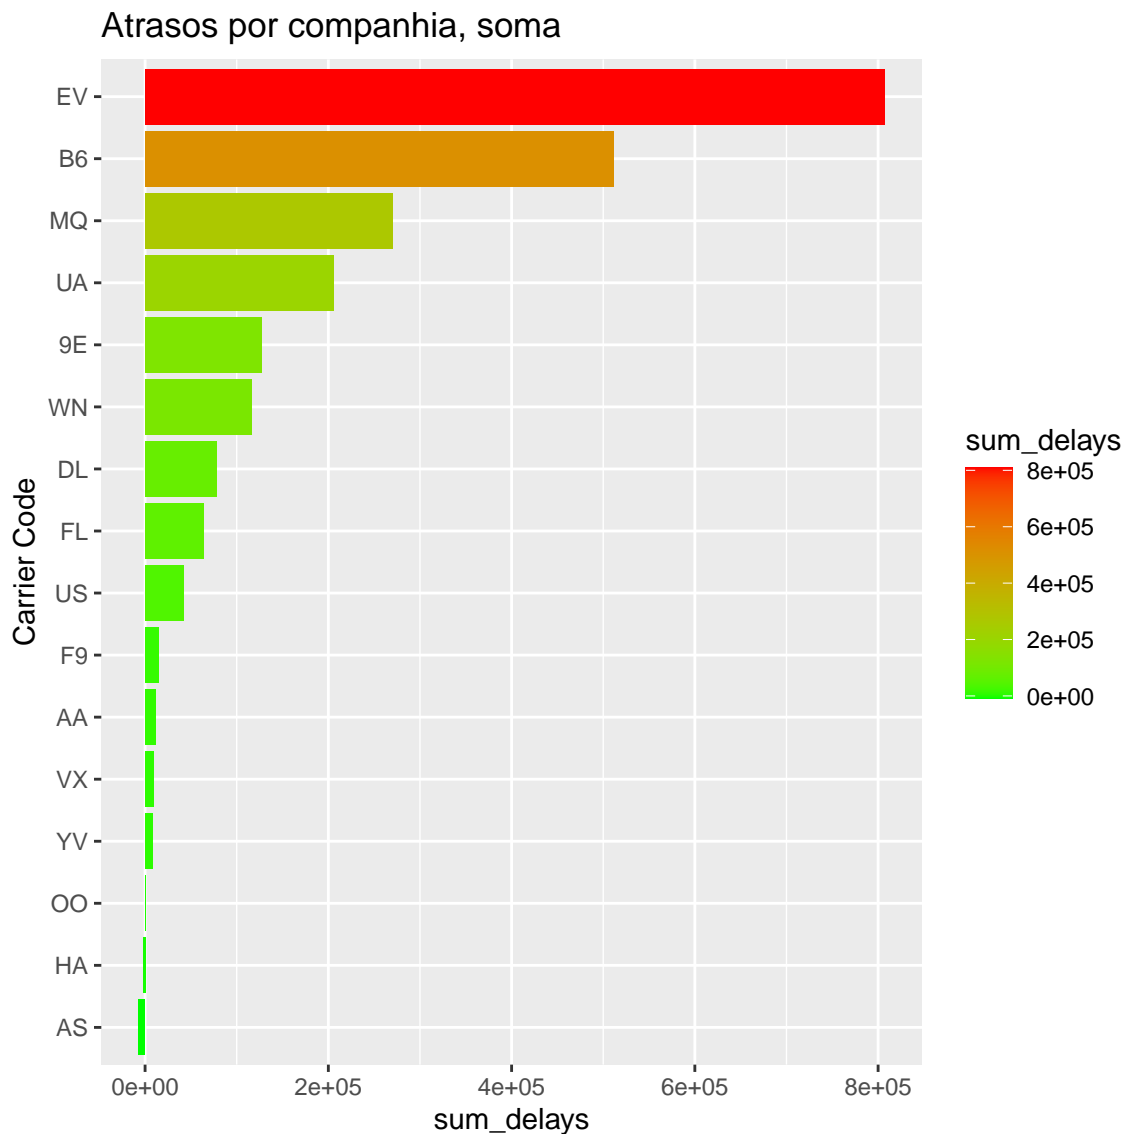
Por mais que realmente existam menos voos marcados durante fins de semana parece que muitos mais são cancelados durante a semana do que em fins de semana.

### 5.6.5

Which carrier has the worst delays? Challenge: can you disentangle the effects of bad airports vs. bad carriers? Why/why not? (Hint: think about flights `%>% group_by(carrier, dest) %>% summarise(n())`)

Para essa Análise me parece o mais razoável analisar somente o `arr_delay`, visto que por mais chato que seja ficar aguardando no aeroporto no final o que mais importa é se o avião chega ao seu destino final no horário, para os passageiros.


```
not_cancelled %>%  
  group_by(carrier) %>%  
  summarise(sum_delays = sum(arr_delay), .groups = 'drop') %>%  
  ggplot(mapping = aes(x = reorder(carrier, X = sum_delays), y = sum_delays, fill = sum_delays)) +  
  geom_col() +  
  coord_flip() +  
  scale_fill_gradient(low="green", high="red") +  
  xlab("Carrier Code") +  
  ggtitle("Atrasos por companhia, soma")
```



Por mais bonito que o gráfico possa ficar o mais justo seria considerar o atraso médio, não o total de atrasos no período (ou ainda vale pensar qual a medida estatística mais adequada, por exemplo uma medida mais robusta poderia ser mais interessante para o passageiro médio).

Feitas essas considerações seguimos com o exercício.

```
not_cancelled %>%
  group_by(origin) %>%
  summarise(sum_delays = sum(arr_delay), .groups = 'drop') %>%
  ggplot(mapping = aes(x = reorder(origin, X = sum_delays), y = sum_delays, fill = sum_delays)) +
  geom_col() +
  coord_flip() +
  scale_fill_gradient(low="green", high="red") +
  xlab("Aeroporto saída") +
  ggtitle("Atrasos por aeroporto, soma")
```



É bastante evidente que as companhias que mais voam de EWR inevitavelmente incorrerão em mais atrasos na média.

Poderíamos tentar um score que ponderasse o atraso médio dos aeroportos na avaliação das companhias aéreas, mas ainda assim é possível argumentar que as responsáveis pelos atrasos nos aeroportos sejam as companhias que lá mais operam.



### 5.6.6

What does the sort argument to count() do. When might you use it?

```
# ?count
```

O argumento sort, se == TRUE fará com que na saída da função count() os maiores grupos estarão no topo, ordenados.

## 5.7 grouped mutates (and filters)

### 5.7.1

Refer back to the lists of useful mutate and filtering functions. Describe how each operation changes when you combine it with grouping.

### 5.7.2

Which plane (tailnum) has the worst on-time record?

### 5.7.3

What time of day should you fly if you want to avoid delays as much as possible?

### 5.7.4

For each destination, compute the total minutes of delay. For each flight, compute the proportion of the total delay for its destination.

### 5.7.5

Delays are typically temporally correlated: even once the problem that caused the initial delay has been resolved, later flights are delayed to allow earlier flights to leave. Using lag(), explore how the delay of a flight is related to the delay of the immediately preceding flight.

### 5.7.6

Look at each destination. Can you find flights that are suspiciously fast? (i.e. flights that represent a potential data entry error). Compute the air time of a flight relative to the shortest flight to that destination. Which flights were most delayed in the air?

### 5.7.7

Find all destinations that are flown by at least two carriers. Use that information to rank the carriers.

### 5.7.8

For each plane, count the number of flights before the first delay of greater than 1 hour.